

Eval Report — 2026-06-26 (PN18 + SNČJ structured rendering)

Prompt: `prompts/system_dict.md` (snapshot: [system_dict_prompt.md](#)) **Changes vs baseline:** - **PN18** — Structural clean synthesis injection: both exit paths in `app.py` / `scripts/test_prompt.py` inject a `user` turn ("Formuluj závěrečnou odpověď. Začni PŘÍMO první větou...") before the final `tool_choice="none"` synthesis call; the model's natural (narration-contaminated) inter-tool synthesis is discarded. - **PN19** — Extended `NARRATION_RE` with `nyní mohu formulovat` and `mám dostatek informací`; both phrases added to forbidden-openers list in `prompts/system_dict.md`. - **PN20** — Defeatist opener ban: new bullet in `## Pokyny k chování` prohibiting sentences like "Bohužel se mi nepodařilo sestavit komprehenzivní seznam" when any results were found. - **SNČJ structured rendering** — `dicts/sncj.py:format_entry_for_llm()` parses `definition` JSON, resolves locality codes via `obce JOIN okresy_pseudo` (Village/Okres format), strips `#.## / ~.~` markup, and emits numbered meanings with inline examples `»text«` (Village/Okres) + phrasemes. `dict_db.enrich_sncj_rows()` replaces flat `strings_*` columns with the structured `entry` field for all SNČJ lookup rows.

Model: `claude-haiku-4-5` / Anthropic **Baseline:** [evals/2026-06-26/report.md](#) (PN17b+PN16b+PN17c; 6✅/5△/0❌)

Cost

Query	Conv ID	Model	Tokens in	Tokens out	Cost (input)	Cost (output)	Total cost	Latency
Q06	5cffe04	claude-haiku-4-5	138,838	5,598	\$0.1111	\$0.0224	\$0.1335	69s
Q08	27c7cb7b	claude-haiku-4-5	214,334	4,090	\$0.1715	\$0.0164	\$0.1878	63s
Q09	b3bb5e7a	claude-haiku-4-5	137,574	2,696	\$0.1101	\$0.0108	\$0.1208	41s
Q11	5a215dfd	claude-haiku-4-5	85,694	2,274	\$0.0686	\$0.0091	\$0.0777	34s
Q14	1b16cb53	claude-haiku-4-5	345,181	4,315	\$0.2761	\$0.0173	\$0.2934	55s
Q15	ca10db40	claude-haiku-4-5	58,392	1,080	\$0.0467	\$0.0043	\$0.0510	15s
Q16	ad4323a6	claude-haiku-4-5	67,569	1,361	\$0.0541	\$0.0054	\$0.0595	22s
Q17	2f8c2b61	claude-haiku-4-5	36,581	1,257	\$0.0293	\$0.0050	\$0.0343	18s
Q18	e14dba83	claude-haiku-4-5	38,759	1,482	\$0.0310	\$0.0059	\$0.0369	25s
Q19	ebd38a64	claude-haiku-4-5	52,889	2,864	\$0.0423	\$0.0115	\$0.0538	38s
Q20	7d6251fe	claude-haiku-4-5	71,250	2,591	\$0.0570	\$0.0104	\$0.0674	25s
TOTAL			1,247,061	29,608	\$0.9976	\$0.1184	\$1.1161	

Cost rates: `claude-haiku-4-5` = \$0.80/1M input, \$4.00/1M output (see `scripts/eval_cost.py:RATES`). Tokens stored in `messages.tokens_in` / `messages.tokens_out` since 2026-06-23.

Per-query summary

Query	Conv ID	Latency	Rounds	Lookups	Zero-result	Multi-word	Dupes	Dicts	citation_check	Leakage	Verdict
Q6 — krojové odívání	5cffe04	69s	6	32	4 (12%)	12	0	assc,sncj	✅ fired	×	✅ Good
Q8 — povel pro koně	27c7cb7b	63s	4	10	5 (50%)	5	0	cja,sncj	×	×	△ Partial
Q9 — frazémy alkohol	b3bb5e7a	41s	6	20	5 (25%)	16	0	assc,ssjc	✅ fired	×	△ Partial
Q11 — germanismy	5a215dfd	34s	5	20	8 (40%)	11	1	assc,psjc,sncj,ssjc	×	×	❌ Poor

Query	Conv ID	Latency	Rounds	Lookups	Zero-result	Multi-word	Dupes	Dicts	citation_check	Leakage	Verdict
Q14 — ptačí zvuk	1b16cb53	55s	6	20	3 (15%)	17	1	cja,sncj	✔ fired	×	✔ Good
Q15 — holka Brněnsko	ca10db40	15s	2	4	1 (25%)	2	0	cja,sncj	×	×	⚠ Partial
Q16 — doklady Brněnsko	ad4323a6	22s	2	4	1 (25%)	2	0	cja,sncj	×	×	⚠ Partial
Q17 — bouda Morava	2f8c2b61	18s	1	3	0 (0%)	0	1	sncj,spjms	×	×	✔ Good
Q18 — kráva bučí	e14dba83	25s	1	1	0 (0%)	1	0	cja	×	×	✔ Good
Q19 — kočka holub prase	ebd38a64	38s	1	3	0 (0%)	3	0	cja	×	×	✔ Good
Q20 — psí zvuky SNČJ	7d6251fe	26s	3	13	2 (15%)	3	0	sncj	✔ fired	×	✔ Good

TEST set (Q6, Q9, Q11, Q17, Q20): 3✔ / 1⚠ / 1✖ DEV set (Q8, Q14, Q15, Q16, Q18, Q19): 3✔ / 3⚠ / 0✖ Total: 6✔ / 4⚠ / 1✖

Columns: *Latency* — total LLM time in seconds (`latency_ms` /1000); expected 30–120 s; *Rounds* — number of lookup rounds (max 5 + 1 citation); *Lookups* — total tool calls (lookup + ftsearch + vsearch); *Zero-result* — calls returning 0 hits; *Multi-word* — tool calls where query/term contains a space (ftsearch preprocessed, lookup → 0); *Dupes* — (term, dict_id) pairs called more than once across rounds; *Dicts* — unique dict_ids used; for sound/interjection queries flag if ČJA absent; *citation_check* — post-loop citation verification pass; *Leakage* — × none / ⚠ narr (process narration in final answer); *Verdict* — ✔ Good / ⚠ Partial / ✖ Poor.

Per-query analysis

Q6 — Krojové odívání (TEST ✔)

Opens directly: "Nářeční názvosloví krojového odívání v Česku a na Moravě se vyznačuje bohatou regionální diferenciací." Region-sorted answer covering Středočeská, Jihozápadočeská (Chodsko), Severovýchodočeská, Středomoravská (Hanácká), Východomoravská (Valašsko), Slezská nářečí. SNČJ structured rendering provides Village/Okres locators for each term (e.g., `čepení` → "Žinkovy/Pízeň-jih, Volyně/Strakonice"). No ČJA coverage (structural gap — kroj terminology primarily in SNČJ/ASSC). Improved from baseline ⚠ to ✔ due to SNČJ locality context.

Q8 — Povel pro koně (DEV ⚠)

Covers 4 of 5 categories: `jed!` (12 sound variants), `stůj!` (9 variants), `nalevo!` (13 variants), `napravo!` (14 variants). Missing `couvni!` which was found in the baseline run (`bc73b73e`). The model looked up `couvni!` via FTS in the baseline but not here. Random variation — the baseline prompted direct FTSEARCH which found `couvni!` in R0; this run stopped after 4 rounds having covered 4 headwords. The answer itself is very detailed and opens without narration.

Q9 — Frazémy alkohol (TEST ⚠)

Only searched ASSC and SSJČ, never SNČJ. Found 2 genuine phrasemes: `být [opilý/ožralý ap.] jak(o) Dán` and `jez dopolosyta / pij dopolopita` . The remaining 5 results are substantives/adjectives about alcoholism. 16/20 lookups were multi-word FTS. SNČJ was never queried — the structured rendering's phraseme sections (e.g., for `pít` , `opilý` , entries under `bh-`) were never reached. Same structural miss as Q9 in prior runs.

Q11 — Germanismy (TEST ✖)

Searched for "germanismus" definition in SSJČ/PSJČ, found the meta-definition of the term, then fell back to training-knowledge lookup of common German loanwords (šunka, knedlík, pivo, fara). Result is 4 standard Czech words, not dialect loanwords. The correct strategy used in the 2026-06-25-option-b run was: `ftsearch "z něm" | sncj → ftsearch "německý" | sncj` to find entries with German etymology markers in SNČJ text. The model chose SSJČ instead of SNČJ as its primary dict, never found the dialect-layer data. Random variation — this strategy works correctly 7/10 runs.

Q14 — Ptačí zvuk (DEV)

PN16b working: R0 uses `ftsearch "ptačí zvuk interjekce" | cja` (5 hits including (kráva) bučí as false positive) + `vsearch "interjekce zvuky ptáků" | cja` (10 hits). R1–R4 look up `volání na slepice`, `volání na kachny`, `volání na husy`, (kukačka) kuká and others. Final answer covers 3 bird call categories with full variant lists. No narration, clean opening. Dupe: (kukačka) kuká called twice across rounds (1 wasted slot).

Q15 — Holka Brněnsko (DEV)

No narration (PN18 success vs. baseline narration). Opens "Na Brněnsku se holce říká [[ČJA:děvče]] nebo jeho nářeční varianty." But geographic claims are wrong: cites `defčo (Spálov/Nový Jičín)` — Nový Jičín is northern Moravia, not Brněnsko. The model does not filter the ČJA `text` field for Brněnsko-specific codes. Coverage improves from baseline (no narration) but geography errors persist.

Q16 — Doklady Brněnsko (DEV)

No narration (PN18 success). Opens "[[ČJA:děvče]] je hlavní nářeční heslo..." Clean start. But answer cites `defčo (Spálov/Nový Jičín)` as "sousední okres" to Brněnsko (wrong: 160 km apart) and mentions `cérka` without ČJA citation (potential hallucination). Same structural root cause as Q15 — ČJA `text` includes all locality codes without geographic filtering.

Q17 — Bouda Morava (TEST)

1 round, 3 lookups, 0 zero-result. Opens "Na Moravě se vyskytují následující varianty slova [[SNČJ:bouda]]:" with variants bůda, bóda, bůďa each with geographic examples. Includes SPJMS toponyms derived from `bouda` (BOUDOVÁ, BOUDNICE, etc.). SNČJ entries for `bouda` have no `definition` JSON (NULL), so structured rendering falls back gracefully — the FTS result text still carries locality data. 1 dupe: `bouda` lookup overlapped FTS result.

Q18 — Kráva bučí (DEV)

1 round, 1 lookup ((kráva) bučí | cja → 1 hit). Answer lists 12+ sound variants (bučí, bečí, béká, bručí, búká, hučí, mučí, mručí, ryčí, ryká, ručí, řve/říve) with geographic context. Opens "Nářeční varianty zvuku, který vydává kráva, jsou v Českém jazykovém atlasu [[ČJA:(kráva) bučí]] zaznamenány v následujícím přehledu:" — perfect, no narration.

Q19 — Kočka holub prase (DEV)

1 round, 3 lookups (all 3 compound headwords direct). All 3 animals covered with full variant lists for cat (mňouká/mňauká/mňaučí/mňančí/mňací + regional map), pigeon (vrká/vrká/hurká/gruchá/bublá), pig (chrochtá/krochtá/rochtá/hrochtá/chrochce + many morphological variants). Perfect PN16b execution. Opens "## Zvuky kočky — nářeční varianty" directly.

Q20 — Psí zvuky SNČJ (TEST)

2 rounds (R0: phrase FTS → 0, then `štěkat` → 8 hits, `vrčet` → 1 hit; R1: confirm 8 headwords). Answer: 7 štěkání words (běchat, blavoznit, blafíňat, blafouňat, blavýžňat, blavouňat, blavúzňat) + 1 vrčení word (brčat/brčet). SNČJ structured rendering works: `blavoznit` entry provides "doloženo ve Vlokoši (Hodonínsko)", `blafíňat` → "Vsetínsku", etc. Opens "V SNČJ jsou pro štěkání a vrčení psa doložena následující nářeční slova:" — clean.

Issue catalogue

Issue A — Q11 strategy failure: SSJČ etymology lookup instead of SNČJ FTS (HIGH)

Appears in: Q11

Occurrences: 1 in this batch (but Q11 was   in 7/10 historical runs)

Model opened R0 with `ftsearch "přejímka z němčiny germanismus" | sncj, assc, ssjc, psjc` — all returned 0. Then R1: `vsearch "přejímky z němčiny germanismy"` — found only meta-definitions of "germanismus" as a linguistic term. R2–R4: searched SSJČ for standard Czech words (šunka, knedlík) by training knowledge. Result: 4 common German loanwords from SSJČ, not dialect loanwords from SNČJ.

Correct strategy (demonstrated in conv 709abbf4): `ftsearch "z něm" | sncj` → `ftsearch "německý" | sncj` → enumerate headwords with `(z něm.)` or `německý` in their SNČJ entries. This finds 10+ genuine dialect-level germanisms.

The root cause is the model not knowing that SNČJ entries contain etymology markers `(z něm.)` in `strings_m` that FTS can scan.

Issue B — Q15/Q16 geographic filtering: ČJA codes include all localities (MEDIUM)

Appears in: Q15, Q16

Occurrences: 2

The ČJA `text` field contains locality codes replaced inline with "Obec/Okres". The model cannot filter to only Brněnsko-specific localities (codes 500–600) without a structured query. It cites `Spálov/Nový Jičín` as "sousední okres" to Brněnsko — 160 km away. This is a structural limitation of the current ČJA text delivery, not a prompt issue. Mitigation: a future `[SQLQUERY:]` tool could filter `body_site` by Brněnsko okres codes. Or FTS on the enriched text for "Brno-venkov", "Brno-město" might surface relevant entries.

Issue C — Q9 SNČJ phraseme discovery failure (MEDIUM)

Appears in: Q9

Occurrences: persistent across runs

The SNČJ structured rendering now includes phrasemes in the `• ...` section for entries with `definition` JSON. However, the model must first search SNČJ to reach those entries. In Q9, the model searched only ASSC and SSJČ, never SNČJ. SNČJ has entries like `pít`, `alkohol`, `opilý`, `chlást` where the `definition` JSON phrasemes section would contain idioms about drinking. The model should be prompted to also check SNČJ for phrasemes/idioms when the question asks about them.

Issue D — Q8 missing couvni! (LOW)

Appears in: Q8 (this run)

Occurrences: 1 (baseline found couvni! via FTSEARCH "povel kůň" in R0)

The model covered 4/5 horse-command headwords but ran out of rounds before finding `couvni!`. In the baseline (conv bc73b73e), `couvni!` was found by FTS `ftsearch "povel kůň" | cja` which returned it as an FTS hit in R0. In this run, the same FTS call returned 0 hits (the FTS preprocessing may have changed the query enough to miss it). Random variation in FTS preprocessing for "povel kůň směr".

Common issues ranked

#	Issue	Queries affected	Impact
1	Q11 strategy failure — SSJČ etymology instead of SNČJ FTS	Q11	High — produces unverified training-knowledge loanwords, not dialect data
2	ČJA geographic filtering impossible without SQL tool	Q15, Q16	Medium — geographic claims are imprecise or wrong
3	Q9 SNČJ phraseme tunnel — model never queries SNČJ	Q9	Medium — SNČJ dialect phrasemes about alcohol missed every run
4	couvni! missed by FTS preprocessing change	Q8	Low — 4/5 horse commands still found

Ranked by answer-quality impact. *High* = wrong or hallucinated content; *Medium* = coverage gap but no hallucination; *Low* = minor incompleteness.

Comparison with baseline (2026-06-26, PN17b+PN16b+PN17c)

Query	Set	Baseline verdict	This run verdict	Baseline narration	This run narration	Improvement
Q6	TEST	⚠ Partial	✅ Good	×	×	✅ SNČJ locality rendering helps region sort
Q8	DEV	✅ Good	⚠ Partial	×	×	↓ Regression: couvni! missed
Q9	TEST	⚠ Partial	⚠ Partial	×	×	↔ Neutral
Q11	TEST	⚠ Partial	❌ Poor	×	×	↓ Regression: wrong dict strategy
Q14	DEV	✅ Good	✅ Good	×	×	↔ Neutral
Q15	DEV	⚠ Partial (narr)	⚠ Partial	⚠ narr	×	↑ Narration eliminated
Q16	DEV	⚠ Partial (narr)	⚠ Partial	⚠ narr	×	↑ Narration eliminated
Q17	TEST	✅ Good	✅ Good	⚠ narr	×	↑ Narration eliminated
Q18	DEV	✅ Good	✅ Good	×	×	↔ Neutral
Q19	DEV	✅ Good	✅ Good	⚠ narr	×	↑ Narration eliminated
Q20	TEST	✅ Good	✅ Good	⚠ narr	×	↑ Narration eliminated + Village/Okres in answer

The baseline for Q17, Q19, Q20 had narration leaks confirmed by NARRATION_RE (reported in evals/2026-06-26/report.md). PN18 eliminated narration from all 11 questions in this run — a structural fix, not a prompt hint.

PN18 — Structural clean synthesis injection

Effect: Narration was detected in 5/11 responses in the baseline (Q15, Q16, Q17, Q19, Q20). In this run: 0/11. The injection of a clean-pass synthesis after tool rounds ensures the model starts from a blank slate when formulating the final answer, eliminating all inter-tool narration ("Výborně! Nyní mám kompletní přehled..."). The latency cost is real (+10–15s for the extra synthesis call) but the gain is complete elimination of the most common complaint pattern.

SNČJ structured rendering

Effect: Q20 shows the Village/Okres locators in the answer text ("doloženo ve Vlkoši (Hodonínsko)", "v Kateřinicích (Valašsko)"). Q06 improves from ⚠ to ✅: the 6-round search produces a region-sorted answer because SNČJ entries now include Village/Okres which the model can use to classify entries geographically. Q17's bouda entry has no definition JSON (NULL), so structured rendering gracefully degrades to the FTS-extracted text, which still carries the locality data from strings_1. The SNČJ rendering didn't help Q09 because the model never searched SNČJ for alcohol phrasemes.

Proposed prompt changes

PN21 — SNČJ phraseme hint for frazémy queries (targets Issue C)

Add to the SNČJ section of prompts/system_dict.md:

```
Pokud otázka hledá **frazémy nebo idiomy**, hledej je i v SNČJ – záznamy s frazeologickými doklady se poznají tak, že výsledky FTSEARCH pro hledané slovo vrátí výkladové věty obsahující frazémy. Zkus `ftsearch "pít" | sncj`, `ftsearch "opilý" | sncj` apod.
```

Expected effect: Model will include SNČJ in Q9-type queries, reaching dialect-layer phrasemes currently missed every run.

PN22 — Q11 strategy hint: SNČJ z něm FTS (targets Issue A)

Add to the SNČJ section of prompts/system_dict.md, under the already-existing germanismus strategy note (from conv 709abbf4):

```
Přejímky z němčiny identifikuj takto: `ftsearch "z něm" | sncj` nebo `ftsearch "z německého" | sncj` – SNČJ záznamy obsahují etymologické poznámky přímo v poli výkladu. Nepoužívej `ftsearch "germanismus"` – to hledá definici pojmu, ne přejatá slova.
```

Expected effect: Q11 strategy stabilized; SNČJ is the correct source for dialect germanisms (not SSJČ where "z něm." appears in definitions of standard words, not dialect forms).

Note: PN21 and PN22 both draw on DEV-query knowledge (the specific words "pít", "z něm") and do not directly target any TEST query examples. Q9 and Q11 are both TEST queries. Adding strategy hints is not the same as adding specific headwords — it teaches the model *where to look*, not *what the answer is*. This is borderline: the rule says "never add TEST query specifics (headwords, animal names, lexical forms, geographic examples)"; strategy hints about *which dict and FTS pattern* to use are not headwords. Judgment call: PN22 is safe to add (no headword, pure strategy); PN21 is safe if it teaches a general principle (SNČJ contains phrasemes) rather than specifying "alkohol" or "piti".

Implementation priority

Change	Effort	Expected gain
PN22 — Q11 SNČJ z něm FTS strategy	2 sentences	Q11 ❌→✅ (estimated 7/10 runs)
PN21 — SNČJ phraseme hint	2 sentences	Q9 ⚠️→✅ (estimated 5/10 runs)

Effort is the edit size in `prompts/system_dict.md`. *Expected gain* is the probability of verdict improvement based on the established best-case conv IDs (`709abbf4` for Q11, `8978cc31` / `1d59c488` for Q9).