

Eval Report — 2026-09-05 (Full batch: qwen3.8-flash, plan + rethink)

Prompt: `prompts/system_dict.md` (unchanged from baseline) **Provider/Model:** qwen3.8-flash via OpenCode.ai ("Console Go" backend) **Feature flags:** `rerank=True`, `hyde_v2=True` in all conditions, plus `plan` or `rethink` per condition **Scope:** all 11 currently-answerable queries × 3 conditions (baseline/plan/rethink) = 33 attempted runs, **26 succeeded** **Related:** [evals/2026-09-05-gemma4-full/report.md](#) — the equivalent batch on `google/gemma-4-31b-it`, run the same day with the same methodology. See the comparison section at the end of this report.

Headline finding: 7/33 runs failed (21%)

Before any quality analysis: qwen3.8-flash via OpenCode.ai failed on over a fifth of attempted runs, in two distinct ways.

3 failures — content-moderation false positive, on Q14 (bird interjections), Q17 (bouda/Morava, rethink), Q20 (dog barking, baseline):

```
litellm.BadRequestError: OpenAIException - Error from provider (Console Go):
Upstream request failed: [data_inspection_failed] <400>
InternalError.Algo.DataInspectionFailed: Input text data may contain
inappropriate content.
```

None of these queries involve remotely sensitive content — this is a false positive from an upstream content-safety filter, not a real policy violation. No documented way to disable it was found in OpenCode.ai's docs. Re-tested the same model via OpenRouter (`qwen/qwen3.8-flash`) to see if a different gateway would avoid the filter — its error attributed the request directly to `"provider_name": "Alibaba"`, meaning both paths route to the same upstream backend and would hit the same filter.

4 failures — silent mid-flow death, on Q11 (all three conditions) and Q09 (rethink): the process produced normal output through the citation-check correction pass, printed `"► Formulují opravenou odpověď..."`, then terminated with **no exception, no traceback, no captured error** — not a timeout pattern, not clustered in time (failures occurred 25+ minutes apart), so not one shared infrastructure event. Root cause not identified.

This failure rate is structural for this workload, not a fluke of one run — Q11 failed in all three conditions independently. Combined with `evals/2026-09-05-gemma4-full/`'s 0/33 failures, this is the single most important data point for choosing a production default: `gemma-4-31b-it` is measurably more reliable for dictpert's dict-mode workload.

Cost

Model is treated as flat-rate under `scripts/eval_cost.py:RATES`'s `qwen` prefix match (consistent with `qwen3.5/3.6-plus`) — **pricing not independently confirmed** for `qwen3.8-flash` specifically; reported as \$0.00 accordingly. Token volume is real and notably higher than `gemma-4-31b-it`'s batch (`qwen3.8-flash` produces longer, more exploratory answers — see rounds/lookups below).

Aggregate results by condition (successful runs only)

Condition	Successful runs	Avg latency	Avg rounds	Avg lookups	Zero-result rate	citation_check fired	Leakage
baseline	8/11	59s	5.1	11.9	14% (13/95)	5/8	0/8
plan	10/11	67s	6.4	14.0	10% (14/140)	8/10	1/10

Query	Cond	Conv ID	Latency	Rounds	Lookups	Zero	Dicts	CiteChk	Cites	Leak
Q11	rethink	—	—	—	—	—	—	—	—	FAILED (silent)
Q14	rethink	52226880	102s	8	34	2	assc,cja,sncj,ssjc	✓	20	×
Q15	rethink	88bcb44a	80s	6	13	1	cja,sncj,ssjc	✓	7	×
Q16	rethink	a0eca597	28s	5	10	1	cja	×	3	×
Q17	rethink	—	—	—	—	—	—	—	—	FAILED (moderation)
Q18	rethink	b8272ab0	29s	2	3	0	cja,sncj	×	2	×
Q19	rethink	eabea0a8	56s	1	3	0	cja	×	3	×
Q20	rethink	b2a73f61	74s	7	23	2	sncj	×	13	×

Comparison with google/gemma-4-31b-it (evals/2026-09-05-gemma4-full/)

Same day, same methodology, same 11 queries × 3 conditions.

	qwen3.8-flash (OpenCode.ai)	google/gemma-4-31b-it (OpenRouter)
Failure rate	7/33 (21%)	0/33 (0%)
Cost	\$0.00 (unconfirmed flat-rate assumption)	\$0.27 confirmed
Avg lookups/query	12-16	7-8
Avg latency	59-69s	52-76s
Zero-result trend	14% → 10% → 5% (baseline→plan→rethink)	8% → 3% → 1%
Leakage	1/26 (instr)	0/33
plan+rethink mechanism	Works, same directional improvement	Works, same directional improvement
Tool-calling compatibility	Required a fix (thinking-mode disable) to work at all	Worked out of the box

Both models validate the plan/rethink mechanism — the zero-result-rate improvement direction holds on both, which is good independent confirmation that the effect isn't specific to one model. But qwen3.8-flash's 21% failure rate is disqualifying for a production default regardless of its other qualities (richer, more exploratory answers; comparable or better zero-result rates when it *does* succeed). A user hitting a 1-in-5 chance of a hard failure — with no retry path built for the "silent death" failure mode specifically — is a worse experience than gemma-4-31b-it's clean run.

Recommendation: google/gemma-4-31b-it for production default. qwen3.8-flash remains interesting as a secondary/experimental option once (if) the silent-failure and moderation-false-positive issues are root-caused, but isn't ready to be a primary path today.